

```

# Downloading the gene_expression.tsv file
download.file("https://raw.githubusercontent.com/ghazkha/Assessment4/refs/heads/main/gene_expression.tsv",
              "gene_expression.tsv")

# Install and load required packages
if (!require("seqinr")) install.packages("seqinr")

## Loading required package: seqinr
# Load seqinr package
library(seqinr)

# Read downloaded file
gene_expression <- read.table("gene_expression.tsv", header = TRUE, sep = "\t", row.names = 1)

# Display the first few rows of the gene expression data
head(gene_expression)

##                                GTEX.1117F.0226.SM.5GZZ7 GTEX.1117F.0426.SM.5EGHI
## ENSG00000223972.5_DDX11L1                                0                      0
## ENSG00000227232.5_WASH7P                                187                     109
## ENSG00000278267.1_MIR6859-1                             0                      0
## ENSG00000243485.5_MIR1302-2HG                            1                      0
## ENSG00000237613.2_FAM138A                               0                      0
## ENSG00000268020.3_OR4G4P                                0                      1
##                                GTEX.1117F.0526.SM.5EGHJ
## ENSG00000223972.5_DDX11L1                                0
## ENSG00000227232.5_WASH7P                                143
## ENSG00000278267.1_MIR6859-1                             1
## ENSG00000243485.5_MIR1302-2HG                            0
## ENSG00000237613.2_FAM138A                               0
## ENSG00000268020.3_OR4G4P                                0

# Create a new column with the mean of the other columns
gene_expression$Mean <- rowMeans(gene_expression)

# Display the first few rows with the new Mean column
head(gene_expression)

##                                GTEX.1117F.0226.SM.5GZZ7 GTEX.1117F.0426.SM.5EGHI
## ENSG00000223972.5_DDX11L1                                0                      0
## ENSG00000227232.5_WASH7P                                187                     109
## ENSG00000278267.1_MIR6859-1                             0                      0
## ENSG00000243485.5_MIR1302-2HG                            1                      0
## ENSG00000237613.2_FAM138A                               0                      0
## ENSG00000268020.3_OR4G4P                                0                      1
##                                GTEX.1117F.0526.SM.5EGHJ          Mean
## ENSG00000223972.5_DDX11L1                                0    0.0000000
## ENSG00000227232.5_WASH7P                                143  146.3333333
## ENSG00000278267.1_MIR6859-1                             1    0.3333333
## ENSG00000243485.5_MIR1302-2HG                            0    0.3333333
## ENSG00000237613.2_FAM138A                               0    0.0000000
## ENSG00000268020.3_OR4G4P                                0    0.3333333

# List the 10 genes with the highest mean expression
top_10_genes <- head(gene_expression[order(-gene_expression$Mean), ], 10)

```

```
# Display the top 10 genes
top_10_genes
```

```
##                                GTEX.1117F.0226.SM.5GZZ7 GTEX.1117F.0426.SM.5EGHI
## ENSG00000198804.2_MT-C01                267250                1101779
## ENSG00000198886.2_MT-ND4                273188                991891
## ENSG00000198938.2_MT-C03                250277                1041376
## ENSG00000198888.2_MT-ND1                243853                772966
## ENSG00000198899.2_MT-ATP6               141374                696715
## ENSG00000198727.2_MT-CYB                127194                638209
## ENSG00000198763.3_MT-ND2               159303                543786
## ENSG00000211445.11_GPX3                464959                39396
## ENSG00000198712.1_MT-C02               128858                545360
## ENSG00000156508.17_EEF1A1              317642                39573
```


```
##                                GTEX.1117F.0526.SM.5EGHJ      Mean
## ENSG00000198804.2_MT-C01                218923 529317.3
## ENSG00000198886.2_MT-ND4                277628 514235.7
## ENSG00000198938.2_MT-C03                223178 504943.7
## ENSG00000198888.2_MT-ND1                194032 403617.0
## ENSG00000198899.2_MT-ATP6               151166 329751.7
## ENSG00000198727.2_MT-CYB                141359 302254.0
## ENSG00000198763.3_MT-ND2               149564 284217.7
## ENSG00000211445.11_GPX3                306070 270141.7
## ENSG00000198712.1_MT-C02               122816 265678.0
## ENSG00000156508.17_EEF1A1              339347 232187.3
```

```
# Count the number of genes with a mean expression less than 10
low_mean_genes <- sum(gene_expression$Mean < 10)
```

```
# Display the number of genes with mean expression < 10
low_mean_genes
```

```
## [1] 35988
```

## Histogram of Gene Expression Means



```
# Downloading the growth_data.csv file
download.file("https://raw.githubusercontent.com/ghazkha/Assessment4/refs/heads/main/growth_data.csv",
              "growth_data.csv")

# Read and inspect the "growth_data.csv" file
growth_data <- read.csv("growth_data.csv")

# Display column names
colnames(growth_data)

## [1] "Site"          "TreeID"        "Circumf_2005_cm" "Circumf_2010_cm"
## [5] "Circumf_2015_cm" "Circumf_2020_cm"


# Calculate the mean and standard deviation of tree circumference at the start (2005) and end (2020) of
circumf_stats <- aggregate(cbind(Circumf_2005_cm, Circumf_2020_cm) ~ Site,
                           data = growth_data,
                           FUN = function(x) c(mean = mean(x), sd = sd(x)))

circumf_stats

##           Site Circumf_2005_cm.mean Circumf_2005_cm.sd Circumf_2020_cm.mean
## 1 northeast          5.2920000         0.9140267          54.22800
## 2 southwest          4.8620000         1.1474710          45.59600
##   Circumf_2020_cm.sd
## 1          25.22795
## 2          17.87345


# Boxplot of tree circumference in 2005 by site
boxplot(Circumf_2005_cm ~ Site, data = growth_data,
```

```
main = "Tree Circumference in 2005 by Site",  
xlab = "Site",  
ylab = "Circumference (cm)")
```



```
# Boxplot of tree circumference in 2020 by site  
boxplot(Circumf_2020_cm ~ Site, data = growth_data,  
        main = "Tree Circumference in 2020 by Site",  
        xlab = "Site",  
        ylab = "Circumference (cm)")
```

## Tree Circumference in 2020 by Site



```
# Calculate mean growth over the last 10 years (2010-2020) at each site
growth_data$Growth_10_years <- growth_data$Circumf_2020_cm - growth_data$Circumf_2010_cm
mean_growth <- aggregate(Growth_10_years ~ Site, data = growth_data, mean)
mean_growth
```

```
##           Site Growth_10_years
## 1 northeast          42.94
## 2 southwest          35.49
```

```
# Perform a t-test to estimate if the 10-year growth differs between the two sites
t_test_growth <- t.test(Growth_10_years ~ Site, data = growth_data)
t_test_growth$p.value
```

```
## [1] 0.06229256
```

## Part 2

```
# Part 2
```

```
# Install required packages
install.packages("R.utils")
```

```
## Installing package into '/home/s224409221/R/x86_64-pc-linux-gnu-library/4.1'
## (as 'lib' is unspecified)
```

```
# Load the necessary libraries
library(R.utils)
```

```

## Loading required package: R.oo
## Loading required package: R.methodsS3
## R.methodsS3 v1.8.2 (2022-06-13 22:00:14 UTC) successfully loaded. See ?R.methodsS3 for help.
## R.oo v1.26.0 (2024-01-24 05:12:50 UTC) successfully loaded. See ?R.oo for help.
##
## Attaching package: 'R.oo'
## The following object is masked from 'package:R.methodsS3':
##
##     throw
## The following object is masked from 'package:seqinr':
##
##     getName
## The following objects are masked from 'package:methods':
##
##     getClasses, getMethods
## The following objects are masked from 'package:base':
##
##     attach, detach, load, save
## R.utils v2.12.3 (2023-11-18 01:00:02 UTC) successfully loaded. See ?R.utils for help.
##
## Attaching package: 'R.utils'
## The following object is masked from 'package:utils':
##
##     timestamp
## The following objects are masked from 'package:base':
##
##     cat, commandArgs, getOption, isOpen, nullfile, parse, warnings
# URL for Saprospirales bacterium CDS
URL_saprospirales <- "https://ftp.ensemblgenomes.ebi.ac.uk/pub/bacteria/release-59/fasta/bacteria_58_col"

# Download the file
download.file(URL_saprospirales, destfile = "saprospirales_cds.fa.gz")

# Read the FASTA file using seqinr's read.fasta
saprospirales_cds <- read.fasta(file = "saprospirales_cds.fa.gz", seqtype = "DNA")
saprospirales_cds_count <- length(saprospirales_cds)

# URL for E. coli CDS
URL_ecoli <- "http://ftp.ensemblgenomes.org/pub/bacteria/release-53/fasta/bacteria_0_collection/escheri"

# Download the file
download.file(URL_ecoli, destfile = "ecoli_cds.fa.gz")

# Read the FASTA file for E. coli using seqinr's read.fasta
ecoli_cds <- read.fasta(file = "ecoli_cds.fa.gz", seqtype = "DNA")

```

```
ecoli_cds_count <- length(ecoli_cds)

# Create a table comparing the number of coding sequences for both organisms
cds_table <- data.frame(
  Organism = c("E. coli", "Saprospirales"),
  Coding-Sequences = c(ecoli_cds_count, saprospirales_cds_count)
)

print(cds_table)

##           Organism Coding-Sequences
## 1           E. coli             4239
## 2 Saprospirales             4527
```

### Difference between two organism:

Saprospirales bacterium has a higher number of coding sequences (4,527) compared to *E. coli* (4,239), indicating a potentially more complex genome and greater functional diversity. This difference suggests that *Saprospirales bacterium* may possess specialized genes that enable it to adapt to specific environmental conditions or ecological niches, whereas *E. coli* has a more streamlined genome optimized for versatility in various laboratory and natural settings.

```
# Calculate the total coding DNA length for both organisms
ecoli_total_coding_dna <- sum(sapply(ecoli_cds, length))
sapro_total_coding_dna <- sum(sapply(saprospirales_cds, length))

# Create a table comparing the total coding DNA lengths
total_coding_table <- data.frame(
  Organism = c("E. coli", "Saprospirales"),
  Total_Coding_DNA = c(ecoli_total_coding_dna, sapro_total_coding_dna)
)

print(total_coding_table)

##           Organism Total_Coding_DNA
## 1           E. coli       3978528
## 2 Saprospirales       4200321
```


### Difference between two organism:

The primary difference between *E. coli* and *Saprospirales bacterium* (GCA\_003448025) lies in the total coding DNA, with *E. coli* having 3,978,528 base pairs and *Saprospirales* having 4,200,321 base pairs. This suggests that *Saprospirales* may possess a slightly larger genome, potentially allowing for more diverse metabolic and environmental adaptations compared to *E. coli*.

```
# Calculate sequence lengths for both organisms
ecoli_cds_lengths <- sapply(ecoli_cds, length)
sapro_cds_lengths <- sapply(saprospirales_cds, length)

# Boxplot of coding sequence lengths
boxplot(ecoli_cds_lengths, sapro_cds_lengths,
  names = c("E. coli", "Saprospirales"),
  main = "Coding Sequence Lengths",
```

```
ylab = "Sequence Length (bp)")
```



```
# Calculate the mean and median of sequence lengths for both organisms
```

```
mean_ecoli <- mean(ecoli_cds_lengths)
```

```
median_ecoli <- median(ecoli_cds_lengths)
```

```
mean_sapros <- mean(sapros_cds_lengths)
```

```
median_sapros <- median(sapros_cds_lengths)
```

```
cat("E. coli: Mean =", mean_ecoli, ", Median =", median_ecoli, "\n")
```

```
## E. coli: Mean = 938.5534 , Median = 831
```

```
cat("Saprospirales: Mean =", mean_sapros, ", Median =", median_sapros, "\n")
```

```
## Saprospirales: Mean = 927.8376 , Median = 690
```

### Difference between two organism:

The comparison of coding sequence lengths between *E. coli* and *Saprospirales bacterium* reveals two key differences. First, while the mean coding sequence length is quite similar between the two organisms, with *E. coli* having a slightly higher mean (938.55 bp) than *Saprospirales* (927.84 bp), the difference lies in their median values. The median length for *E. coli* is notably higher at 831 bp compared to *Saprospirales*, which has a median of 690 bp. This suggests that although the average coding sequence lengths are similar, *Saprospirales* exhibits a wider variation, likely containing more shorter sequences compared to *E. coli*, where the lengths are more evenly distributed around the mean.

```
# Define the standard nucleotides
```

```
nucleotides <- c("A", "C", "G", "T")
```



```

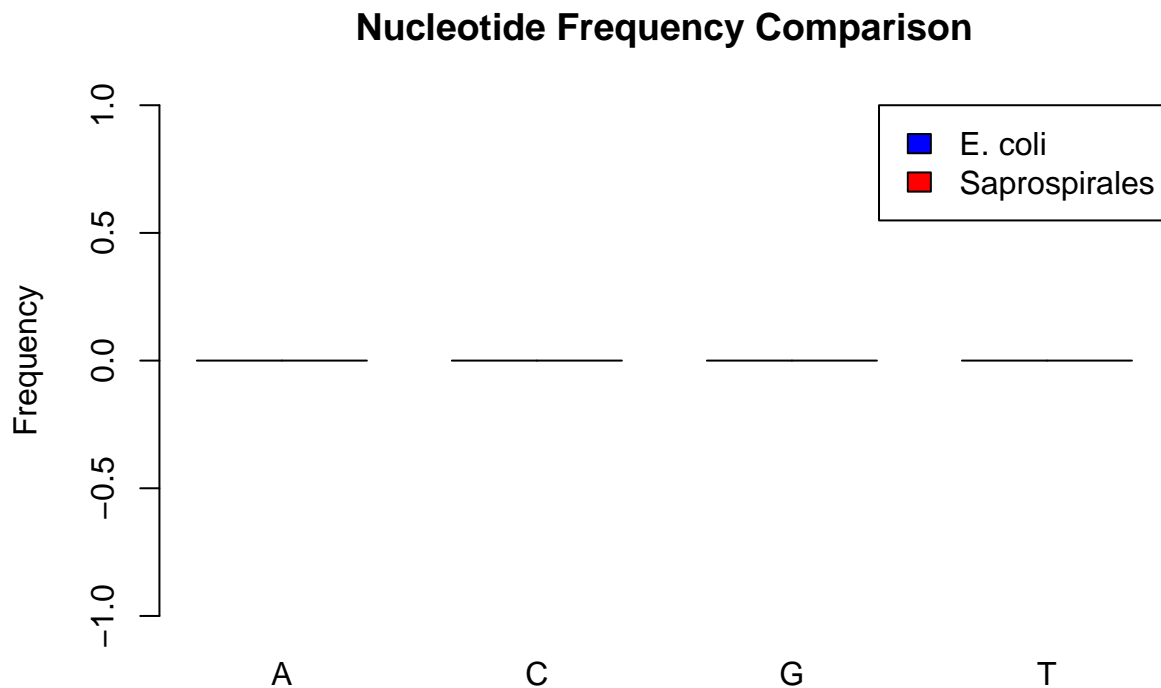
# Combine all E. coli sequences for nucleotide frequency calculation
ecoli_concat <- paste(unlist(ecoli_cds), collapse = "") # Concatenate all sequences
ecoli_nuc_freq <- table(factor(strsplit(ecoli_concat, split = "")[[1]], levels = nucleotides))

# Combine all Saprospirales sequences for nucleotide frequency calculation
sapro_concat <- paste(unlist(saprospirales_cds), collapse = "") # Concatenate all sequences
sapro_nuc_freq <- table(factor(strsplit(sapro_concat, split = "")[[1]], levels = nucleotides))

# Create a matrix for the bar plot
nuc_freq_matrix <- rbind(as.numeric(ecoli_nuc_freq), as.numeric(sapro_nuc_freq))

# Create bar plots for nucleotide frequencies
barplot(nuc_freq_matrix, beside = TRUE,
        names.arg = nucleotides,
        legend.text = c("E. coli", "Saprospirales"),
        main = "Nucleotide Frequency Comparison",
        col = c("blue", "red"),
        ylab = "Frequency",
        args.legend = list(x = "topright"))

```



```

# Calculate codon usage for E. coli
ecoli_codon_usage <- uco(unlist(ecoli_cds), frame = 0, index = "freq")

# Calculate codon usage for Saprospirales
sapro_codon_usage <- uco(unlist(saprospirales_cds), frame = 0, index = "freq")

```

```

# Check the structure of codon usage outputs
str(ecoli_codon_usage)

## 'table' num [1:64(1d)] 0.03362 0.02146 0.0102 0.01738 0.00687 ...
## - attr(*, "dimnames")=List of 1
## ..$ : chr [1:64] "aaa" "aac" "aag" "aat" ...

str(sapros_codon_usage)

## 'table' num [1:64(1d)] 0.0361 0.0224 0.0145 0.0281 0.0113 ...
## - attr(*, "dimnames")=List of 1
## ..$ : chr [1:64] "aaa" "aac" "aag" "aat" ...

# Extract the codon usage frequency values directly (uco might already be a vector)
ecoli_codon_usage_vector <- as.vector(ecoli_codon_usage)
sapros_codon_usage_vector <- as.vector(sapros_codon_usage)


# Ensure both vectors have the same length for comparison
if (length(ecoli_codon_usage_vector) == length(sapros_codon_usage_vector)) {

  # Create a matrix for codon usage comparison
  codon_usage_matrix <- rbind(ecoli_codon_usage_vector, sapros_codon_usage_vector)

  # Create a bar plot to compare codon usage bias
  barplot(codon_usage_matrix, beside = TRUE,
    main = "Codon Usage Bias",
    col = c("blue", "red"),
    ylab = "Codon Usage Frequency",
    legend.text = c("E. coli", "Saprospirales"),
    args.legend = list(x = "topright"))

} else {
  warning("Codon usage vectors for E. coli and Saprospirales have different lengths.")
}

```



### Difference between two organism:

The codon usage bias comparison between *E. coli* and *Saprospirales bacterium* reveals notable differences, as shown in the bar plot. *E. coli* exhibits a more pronounced bias for certain codons, with several codons having visibly higher frequencies (e.g., the tall blue bars), indicating a strong preference for specific codons in its coding sequences. In contrast, *Saprospirales* demonstrates a relatively balanced codon usage distribution, with fewer extreme values, as indicated by the more evenly distributed red bars.

These differences suggest that *E. coli* may have evolved to preferentially use certain codons, potentially due to its tRNA availability or other factors related to translation efficiency. Meanwhile, *Saprospirales* appears to have a less biased codon usage, which could imply different selective pressures or a more diverse set of tRNA genes. This variation in codon usage bias could reflect differences in the organisms' evolutionary history, genome composition, and translational mechanisms.